

Heacheara Rationn



Docs

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations ~

Help ✓

Configuration

Overview [©]

The Entity Builder topology is composed of steps that describe what should be done with the data. Below is a reference of the different steps that can be used to join, filter and transform data in entity builder. Steps are processed asynchronously. Entity Builder uses internal topics to ensure correct relationships.

Topology configuration is provided to Entity Builder in the join.yaml. Topologies must be targeted to building a single output topic.

Step Types

This is a full reference of available topology step configuration. Often times it can be unclear how best to translate your data relationships into an optimal set of steps. For this purpose we have built a topology generator that is designed to be more straight forward to configure and will provide an optimal topology.

Shared Key (=)

The shared key step type (type: shared-key) is a combination of joining data by a shared key and a transformation of the data using a mapper. It requires that all topics being joined contain all of the configured shared key fields in the message key (except in the case of the Join By Value variation). All other fields in the key must be included in additional keys. The shared key fields are read from the message key to build a partitioning key for grouping data for the same key. The fields are combined in the order they are configured. Changing the order of fields in the step configuration will require a reload of all data.

Diagram



About
About
About
Getting Started
Cost
Tour
Architecture

Topology Configuration

Topology Modification

Optimization Guide

Containers

Topology

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

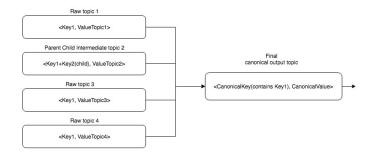
Beam

Getting Started

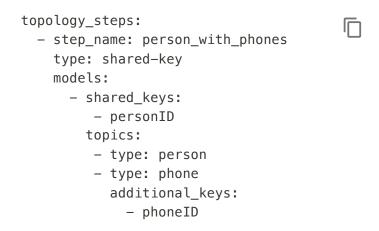
Deployments \

Beam Operations

Help ✓



Example Step Configuration



A full list of configuration options can be found in Shared Key Configuration

How it works

- 1. Entity Joiner receives a message from a topic
- Message is persisted to the persistence table. If the value of the message is null, then the persisted data row is deleted for the message instead.
- 3. A message is produced to the publish queue for the shared key.
- Entity Transformer reads the publish queue message, reads data from the persistence table and sends it to the mapper.
- Mapper returns a canonical message and the Transformer produces to the output topic.

Note: Shared key is the only type that produces messages to the publish queue. A shared key step is required in every topology to create output messages.

Notable Variations

About
About
About
Getting Started
Cost
Tour
Architecture
Topology
Topology Configuration

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help

There are several configurations that vary the functionality of the shared key step.

Join By Value

A shared key step can be configured to join a topic using fields from the value. This requires an additional rekey lookup table to maintain the relationship between the message key and the shared key built using message fields. This additional table is needed to handle clean up when a message with a null value is received or when the field value being joined by changes

Example Step Configuration

You indicate one or more of the shared keys comes from the value of the message by setting

```
from_value on shared_key_overrides.
```

How it works

Create/Update

- 1. Entity Joiner receives a message from a topic
- 2. Reads rekey lookup table to get previous partitioning key
- 3. If previous partitioning key is different from the new partitioning key
 - i. Row for the previous partitioning key is deleted from the persistence table



- About
 About
 About
 Getting Started
 Cost
 Tour
 Architecture
 Topology
 - **Topology Configuration**
 - **Topology Modification**
 - Optimization Guide
 - Containers
 - Database
 - **Test Harness**
 - Template Mapper
 - Transformer
 - Deployments
 - Help

Beam

- Terms of Service
- Getting Started
 - •
 - Deployments
 - Beam Operations
 - Help

- ii. A message is produced to the publish queue for the previous partitioning key
- iii. rekey lookup table row is persisted for the new partitioning key
- 4. if a previous partitioning key was not found
 - i. Row is written to rekey lookup table for the new partitioning key.
- 5. Row is persisted to the persistence table
- 6. A message is produced to the publish queue for the shared key.
- 7. Entity Transformer reads the publish queue message, reads data from the persistence table and sends it to the mapper.
- 8. Mapper returns a canonical message and the Transformer produces to the output topic.

Delete

- 1. Entity Joiner receives a message from a topic
- 2. Reads rekey lookup table to get previous partitioning key
- 3. If previous partitioning key exists
 - i. Row for the previous partitioning key is deleted from the persistence table
 - ii. A message is produced to the publish queue for the previous partitioning key
 - iii. Row is deleted from the rekey lookup table
 - iv. Entity Transformer reads the publish queue message, reads data from the persistence table and sends it to the mapper.
 - v. Mapper returns a canonical message and the Transformer produces to the output topic.

Join Only 👄

^

The <code>join_only</code> option changes shared key step so that rather than producing to the publish queue it produces to the join queue. Data is not sent to the Entity Transformer for transformation, instead it is persisted to be used in later topology steps. It is an intermediate topology step used to join two or more that have a common key, but that key is not the key desired for final output topic. This option offers the benefit of not needing an intermediate data

Entity Builder About About Getting Started Cost Tour Architecture Topology ^

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

transformation step. This step is typically followed by a parent child step or a shared key join by value step

NOTE If the data for a single shared key could exceed 1MB you cannot use the join_only flag. This is due to the way the data is stored in downstream steps. You should instead use a standard shared-key step with a mapper that will produce a message smaller than 1MB with the data you need in downstream steps.

NOTE If not all the data is needed to build the final output message, for instance you need only a single active phone number, then you should consider using a standard shared key and transform the data into an intermediate output message that can be used in later steps.

Overview

Step Types

Shared

Parent

Child

Rekey

Filter

Lookup

Only

Pass

Through

Fan Out

Internal

Topics

Join

Queue

Publish

Queue

Priority

Queue

Priority

Publish

Queue

Configuration

Reference

Low

Join

Low

Transform

Key

Example Step Configuration

```
name: teamCanonical
topology_steps:
  - step_name: person_w_phones
    type: shared-key
    join_only: true
    models:
      - shared_keys:
         - personID
        topics:
         - type: person
         - type: phone
           additional_keys:
             - phoneID
           shared_key_overrides:
             shared key: personID
               override_key: ownerID
               from_value: true
  - step_name: team_sk
    type: shared-key
    models:
      - shared_keys:
          - teamID
        topics:
          - type: team
          - type: person_w_phones
            shared_key_overrides:
```

- shared_key: teamID

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

override_key: teamID>string
 row_type: person
 from_value: true
additional_keys:
 - personID

How it works

This step works identically to the variation of shared key step that is configured without this option. It differs only in that it produces to the join queue instead of the publish queue. **Note** As this flag defines where to produce messages, this option is not allowed if persist only is also set on the step.

Changed Field Tracking

Changed field tracking is a boolean field that creates a change_field_tracking element which entails the version (how many times this canonical schema has been changed) and the changed fields which details the path to the value and how that value was changed ("create","delete","update") which is obtained from the difference from the previously produced canonical. The comparing process is all done automatically in the transformer and all that needs to be done is to turn this feature on.

Example Result

warning In you deploy.yaml you will need to set min.compaction.lag.ms to ensure no cases where canonical updates override each other in kafka before

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations >

Help ∨

they can be compared. This can cause records with relevant changes to be ignored as they will not be considered.

NOTE Changed Field Tracking is not suitable for all cases. Some potential use cases include highly utilized api's with many consumers (allowing consumers to have a higher level of granulairty with what updates they want to include) or canonicals that are primarily consumed into databases that have complexity issues with many updates (and stand to benefit from the potential exclusion of a wide variety of canonical updates)

Example Step Configuration

```
name: teamCanonical
                                         \Box
topology_steps:
  - step_name: person_w_phones
    type: shared-key
    join_only: true
    models:
      - shared_keys:
         - personID
        topics:
         - type: person
         - type: phone
           additional_keys:
             - phoneID
           shared_key_overrides:
             - shared_key: personID
               override_key: ownerID
               from value: true
  - step_name: team_sk
    type: shared-key
    change_field_tracking: true
    models:
      - shared_keys:
          - teamID
        topics:
          - type: team
          - type: person_w_phones
            shared_key_overrides:
              - shared key: teamID
                override_key: teamID>string
                row_type: person
                from value: true
            additional_keys:
```

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations ~

Help ✓

How it works

Change field tracking is an easy-to-use feature with a lot of work being done behind the scenes. When it is active, entity-builder will store canonical records in the persistence store. Then, after an update is processed, the stored record is compared to the new one. This process is done with generic diffing and will populate the value of changedFields on the canonical message. It also affects the amount of processing power the transformer requires.

Changed Field Tracking Array Comparisons

By default changed field tracking will only indicate and compare based on the indices of the elements that have changed in an array. The results can be misleading if the mapper produces messages in an inconsistent order or an element is removed from an array or inserted in the middle of an array. Optionally, for arrays of complex elements, you can configure the fields that make the elements of an array unique. When provided, array element will be compared to the element that has the same unique key. This is done by configuring a list of array comparisions in changed_field_tracking_settings on the sharedkey step. One array comparison config can be provided per array in the canonical schema. It consists of two fields path (the location of the array field the configuration applies to) and an optional list element key fields (the list of fields within each array element that makes the element unique). For primitive arrays only the path is required and the values will be used as the key.

The output for create and update <code>operationTypes</code> remains the same without the array comparison configuration, although the comparison will be made to the previous output that relates to its given key rather the indice in the array. For delete, <code>deletedArrayElements</code> will include the values for the configured <code>element_key_fields</code> that has been removed.

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations >

Help ∨

Example Result

```
"changeTracking": {
    "com.optum.exts.common.ChangedFieldsTrac
      "version": 0,
      "changedFields": {
        "array": [
            "path": "publications[0]",
            "operationType": "create"
            "path": "publications[]",
            "operationType": "delete",
            "deletedArrayElements": {
              "array": [
                "path": "publisher",
                "value": "Me"
            }
          }
        ]
      }
    }
  }
```

Example Step Configuration

```
- step_name: book
                                        type: shared-key
  changed_field_tracking: true
  changed_field_tracking_settings:
    array_comparisons:
      - path: publications
        element_key_fields:
          - publisher
 models:
  - shared_keys:
      -id
    topics:
      - type: book_w_author
        additional_keys:
          - authorID
      - type: book_w_category
        additional_keys:
          - categoryID
      - type: publication
```

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help ✓

shared_key_overrides:
 - shared_key: id
 override_key: bookID
additional_keys:
 - id

Row Type Override

When row types might have potential collisions row_type_override can be used to rename the referenced name of a given row. The example below shows the row type person being overridden. This feature can be used on the child and parent.

Example Step Configuration

```
- step_name: person_w_phones
                                         \Box
    type: shared-key
    join_only: true
    models:
      - shared_keys:
         - personID
        topics:
         - type: person
           row_type_override: prsn
         - type: phone
           additional_keys:
             phoneID
           shared_key_overrides:
             - shared_key: personID
               override_key: ownerID
               from_value: true
```

How it works

Will change the referenced row type to the new overridden name. This will make it so future references to this row should be through the row_type_override name specified. For example references to the row type in future parts of the topology and in the rows received by the mapper.

Multi Model

Shared key steps can define multiple models. This allows different shared keys and different sets of topics to be defined. This variation can be used when two or more distinct sets of topics should be

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help 🗸

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help ✓

transformed into the same output message format and topic.

Example Step Configuration

```
- step_name: coordination_of_benefits
    type: shared-key
    models:
      - shared_keys:
          - PARTN NBR
          - CNSM_ID
          - SRC_CD
          - LGCY SRC ID
          - COB_COV_TYP_CD
          - OI_TYP_CD
          - COB_PRIMACY_EFF_DT
        topics:
          - type: cnsm_cob_prmcy_w_srch_and_
      - shared_keys:
          - PARTN NBR
          - CNSM_ID
          SRC_CD
          - LGCY_SRC_ID
          - MDCR_PART_TYP_CD
          - MDCR_PRIMACY_EFF_DT
        topics:
          - type: cnsm_mdcr_prmcy_w_srch_oth
            additional_keys:
              – COB_COV_TYP_CD
```

How it works

This step works the same as if two shared key steps were configured except that it shares a persistence table.

NOTE The mapper for this step will need to be able to handle data from both sets of topics to produce the correct output messages.

Parent Child

The parent child step type (type: parent-child) allows connecting two topics where there is a relationship between two topics and the key(s) that connect the two topics is in one of the topics' message value. A common use case for this step is joining topics on a foreign key relationship.

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments

Beam Operations ~

Help ✓

The parent-child step produces output messages to the join queue that are used in later steps. The key for the message is the combination of the parent message key and child message key

Example Step Configuration

```
topology_steps:
    - step_name: city_w_state
    type: parent-child
    parent_child_configs:
        child_config:
        type: city
        keys:
        - cityID
    parent_config:
        type: state
        keys:
        - stateID
```

A full list of configuration options can be found in Parent Child Configuration

How it works

Child message

Create/Update

- Entity Joiner receives a message from the child topic
- 2. Reads child table to read previous parent key
- 3. If previous parent key is different from the new parent key
 - i. Row for the previous parent key is deleted from the persistence table
 - ii. A message is produced to the join queue for the previous parent key / child key
 - iii. child lookup table row is persisted for the new partitioning key
- 4. if a previous parent key was not found in the child table
 - i. Row is written to child table for the new parent key
- 5. Row is persisted to the persistence table
- 6. A message is produced to the join queue for the new parent key / child key

Entity Builder About About About Getting Started Cost Tour Architecture Topology Topology Configuration Topology Modification

Optimization Guide
Containers
Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

Delete

- 1. Entity Joiner receives a message from a topic
- 2. Reads child table to get previous parent key
- 3. If previous parent key exists
 - i. Row for the previous parent key is deleted from the persistence table
 - ii. A message is produced to the join queue for the previous parent key / child key
 - iii. Row is deleted from the child table

Parent message

Create/Update

- 1. Entity Joiner receives a message from the parent topic
- 2. Row is written to persistence table
- RowPKs (child keys) are read from persistence table
- 4. A message is produced to the join queue for each child key

Note as child key is needed for the message to the join queue, if there are no child records in the persistence table then no messages will be published.

Delete

- 1. Entity Joiner receives a message from a topic
- 2. Row is deleted from the persistence table
- 3. RowPKs (child keys) are read from persistence table
- 4. A message is produced to the join queue for each child key

Notable Variations

Persist Parent Only

A less common, but still useful in the right circumstances, variation of a parent-child step is persist parent only. In this case both child and parent topics are configured and consumed, but only the parent record data is persisted. For child records a dummy row with an empty value field is persisted.

Docs

Entity Builder About About Getting Started Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database Test Harness Template Mapper Transformer **Deployments** Help Terms of Service Beam ^ **Getting Started** Deployments

Beam Operations

Help

This reduces data duplication and most importantly can even out partition size in the data store. As child records are stored on the parent records partition, if there are thousands of children this can result in uneven partitions in the persistence store.

Note This should not be used to deal with low cardinality lookup table data

*Example Step Configuration

How it works

This works the same as a standard parent-child table except that child records are not stored in the value column. If child data is needed for mapping then the topic should be included in a later shared-key step.

From Key

In some cases you want the value that is pulled and compared to come from the key of the records linked. Some direct examples are when you use a persist-parent-only parent-child step and wish to link it to another parent-child step using the child key. Another potential case is when you want to simplify the process when your linking value is inside another value such as DATA>key_value.

Note In given example book_w_author is a persistparent-only parent-child step with a child key being referenced

*Example Step Configuration

Entity Builder

About

About

About

Getting Started

Cost

Tour

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Database

Containers

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations >

Help ✓

```
- step_name: book_w_author_and_category F
    type: parent-child
    parent child configs:
      child_config:
        type: book_w_author
        keys:
          - authorID
          -id
          - categoryID
        parent_keys:
          - parent_key: id
            child_key: categoryID
            from key: true
      parent_config:
        type: category
        keys:
          - id
```

How it works

This check simply pulls from the key instead of the value when specified.

Low Priority Parent updates

As parent updates trigger reprocessing of all related child keys, a parent update can lead to spikes in lag on the join-queue that can lead to delays in processing for other messages. A parent-child step can be marked as low priority which will divert messages triggered by parent updates to the low priority join queue. Processing of the messages on the low priority join queue is deferred until higher priority messages have been processed.

Note Downstream processing triggered by low priority parent updates will maintain low priority. They will produce to low priority join queue and low priority publish queue

Example Step Configuration

```
- step_name: application_w_category
    type: parent-child
    parent_child_configs:
        low_priority: true
    child_config:
```

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help \

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

How it works

This works the same as a standard parent-child table except that parent updates produce messages for each related child to the low priority join queue to be processed after high priority messages are completed.

Row Type Override

When row types might have potential collisions row_type_override can be used to rename the referenced name of a given row. The example below shows the row type person being overridden. This feature can be used on the child and parent.

Example Step Configuration

```
- step_name: phone_w_person
                                         \Box
    type: parent-child
    parent_child_configs:
      child_config:
        type: phone
        keys:
          - phoneID
        parent_keys:
          parent key: personID
            child_key: ownerID
            output_key: personID
      parent_config:
        type: person
        row_type_override: prsn
        keys:
          personID
```

How it works

Will change the referenced row type to the new overridden name. This will make it so future references to this row should be through the row_type_override name specified. For example

Beam

Help

Getting Started

Beam Operations

Deployments

Docs

Entity Builder About About Getting Started Cost Tour Architecture Topology ^ **Topology Configuration Topology Modification** Optimization Guide Containers Database **Test Harness** Template Mapper Transformer **Deployments** Help Terms of Service

references to the row type in future parts of the topology and in the rows received by the mapper.

Join by a subset of parent keys

Parent child steps can also be configured to join children to parents using only a subset of the parent keys. This will result in many children joined to many parents. This is configured by setting keys that are not joined to the child type in the additional_keys field of parent_config.

 \Box

*Example Step Configuration

```
- step_name: book_by_author
    type: parent-child
    parent_child_configs:
      child_config:
        type: book_author
        keys:
          bookID
          - authorID
        parent_keys:
          - parent_key: bookID
            child_key: bookID
            output_key: bookID
      parent_config:
        type: book
        keys:
          bookID
        additional_keys:
          - edition
```

How it works

^

This works the same as a standard parent-child step except that multiple parent rows will be joined to each child, meaning downstream steps will also receive all parents.

Rekey

The rekey step allows you to add fields from the message value to the message key for joining purposes. This was previously achieved with a parent-child step with a parent type that was not configured in topics.yaml.

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture

Topology Configuration

Topology Modification
Optimization Guide

Containers

Topology

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help ∨

Example Step Configuration

```
topology_steps:
    - step_name: city_by_state_id
    type: rekey
    rekey_config:
        type: city
        keys:
        - cityID
        keys_from_value:
        - path: stateID
```

A full list of configuration options can be found in Rekey Configuration

How it works

Create/Update

- Entity Joiner receives a message from the child topic and builds
- Reads child table to read previous value keys for the message
- 3. If previous value keys are different from the new value keys
 - i. A delete sent to downstream steps with keys_from_value added to the message key based on values in the child table
 - ii. child lookup table row is persisted for the new value keys
- 4. if a previous key was not found in the child table
 - i. Row is written to child table for the new keys_from_value
- Message is passed along to downstream steps with the new keys_from_value added to the message key.

Delete

- 1. Entity Joiner receives a message from a topic
- 2. Reads child table to get previous value keys
- 3. If previous value keys exist
 - i. Key is modified based on values from child table and delete is sent to downstream steps
 - ii. Once downstream steps have all succeeded, row is deleted from the child table

About
About
About
Getting Started
Cost
Tour
Architecture
Topology
Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

Filter ⁽⁻⁾

The filter step type (type: filter) allows a topic to be filtered to remove messages that should not be processed in a subsequent step of the topology.

Example Step Configuration

```
topology_steps:
                                         \Box
  - step_name: phone_filter
    type: filter
    filter_config:
      type: phone
      output_type_prefix: filtered
      any match:
        - conditions:
            - field_name: active
              from_value: true
              blacklist:
                - "false"
            - field_name: phoneType
              whitelist:
                - cell
                home
```

Filter steps can be configured with one to many matches in the any_match field. If any of these matches succeeds, the message will be passed along to downstream steps. Otherwise, the message will be filtered out. Steps that wish to use the filter messages should be configured to use the name of the filter step. Any step that uses the raw topic type will get the unfiltered messages.

Each match can be configured with one to many conditions. All conditions for a match must pass for a match to succeed. A condition can be defined for a field in the message key or message value using the from_value property.

Types of Conditions

Several types of conditions can be configured in a filter step.

Blacklist/Whitelist Condition =



About
About
About
Getting Started
Cost
Tour
Architecture
Topology
Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

A blacklist/whitelist condition can define either a whitelist (the field must match one of the values in the list) or a blacklist (the field must not match one of the values in the list). Configuring "null" in the whitelist or blacklist will match null values for the field.

Examples of Blacklist/Whitelist Condition Configuration

```
- step_name: categorized_books
  type: filter
  filter_config:
    output_type_prefix: categorized
    type: book_data
    any_match:
    - conditions:
    - field_name: categoryID
        from_value: true
        blacklist:
        - "-1"
```

In this example categoryID values with -1 would be filtered out for not meeting the requirements set.

A blacklist/whitelist condition can also specify a wildcard match field which is a boolean value to specify the use of asterisks as a wildcard that wil match 0 or more characters.

In this example categoryID values that match regular expression values "-" or "1" would be filtered out for not meeting the requirements set.



About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

A changed fields condition is used to filter messages so that only changes to configured fields trigger updates.

Example Changed Fields Condition Configuration

```
- step_name: book_title
  type: filter
  filter_config:
    type: book_data
    output_type_prefix: filtered
    any_match:
    - conditions:
        - changed_fields:
        - title
```

In this example only changes made to the configured field title will trigger updates.

Ignore Changed Fields Condition

An ignore changed fields condition is used to filter messages so that changes to listed fields will not trigger updates, while all other changes to other fields will.

Example Ignore Changed Fields Condition Configuration

```
- step_name: book_title
  type: filter
  filter_config:
    type: book_data
    output_type_prefix: filtered
    any_match:
    - conditions:
    - ignore_changed_fields:
    - title
    - authorID
```

In this example any changes to *only* the fields title and/or authorID will be ignored. Changes to fields not in the list will *always* trigger updates.

Please note that you can only have a single ignore changed fields condition in a step, and you also cannot have both an ignore changed fields condition and a changed fields condition in a single step.

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help

Ignoring fields in an array is not supported. You must specify avro schema defaults for ignored fields.

Field Length Condition

A field length condition defines what the length an input field string must be using less than and/or greater than integer configuration values.

Example Field length Condition Configuration

```
- step_name: state_cleaned
  type: filter
  filter_config:
    type: state
    output_type_prefix: filtered
    any_match:
    - conditions:
    - field_name: region
        from_value: true
        field_length:
        less_than: 8
```

In this example all state's region value that have a string length 8 or greater would be filtered out for not meeting the requirements set.

Numeric Condition —

A numeric condition defines numeric boundaries on a numeric input field using less than/less than or equal to and/or greater than/greater than or equal to floating point configuration values.

Example Numeric Condition Configuration

```
- step_name: book_price
  type: filter
  filter_config:
    type: book_data
    output_type_prefix: filtered
    any_match:
    - conditions:
    - field_name: price
        row_type: book
        from_value: true
        numeric_value:
        less_than: 8.50
```

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations ~

Help

```
greater_than_or_equal_to: 1.25
```

In this example all book prices that are 8.50 or greater or less than 1.25 would be filtered out for not meeting the requirements set.

Date Condition (=)

A date condition defines date/timestamp boundaries on a date formatted string input field using before and/or after date/timestamp values. Before and after date configurations are not inclusive and must be in the format of yyyy-mm-dd or yyyy-mm-ddThh:mm:ss.sssZ.

Example Date Condition Configuration

```
- step_name: book_publish_date
  type: filter
  filter_config:
    type: book_data
    output_type_prefix: filtered
    any_match:
    - conditions:
    - field_name: publishDate
        row_type: book
        from_value: true
        date:
        before: 2019-01-01
        after: 2002-01-01
```

In this example all book publish dates that are after January 1st 2019 or before January 1st 2002 would be filtered out for not meeting the requirements set.

The results of a filter step are used in downstream steps by configuring a type of

```
output_type_prefix _ type , not the filter step name. For the example above you would configure filtered_phone
```

A full list of configuration options can be found in Filter Configuration

How it works

Message not filtered



- About
 About
 About
 Getting Started
 Cost
 Tour
 Architecture
 - **Topology Configuration**
 - **Topology Modification**
 - **Optimization Guide**
 - Containers

Topology

- Database
- **Test Harness**
- Template Mapper
- Transformer
- Deployments
- Help
- Terms of Service
- Beam
 - Getting Started
 - Deployments >

^

- Beam Operations
- Help

- 1. Entity Joiner receives a message from a topic
- Configured matches are applied in order. As soon as one match is found that succeeds, the messages is passed through to the downstream steps.
- 3. Message key is stored in the filter table. This keeps track of messages that have passed the filter so if the message is filtered by value it can be cleaned up properly.

Filtered by key field condition

- 1. Entity Joiner receives a message from a topic
- Configured matches are applied in order. No matches pass and at least one condition configured for a key field fails.
- 3. Offset for the message is committed and joiner moves on to the next message.

Note As Kafka message keys are immutable, filtering based on a key field takes precedence over filtering based on a value field. This is because filtering by key is more efficient than by value. No interactions are needed with the filter table and the message is not sent to downstream steps.

Filtered by value field condition

- 1. Entity Joiner receives a message from a topic
- Configured matches are applied in order. No matches pass and all failed conditions are for value fields.
- The filter table is queried for the message key. If the key is not found in the table, the offset for the message is committed, and it is not sent to downstream steps.
- If the key is in the filter table the message is sent to downstream steps and treated as a delete message (null message value).
- Once all downstream steps have cleaned up data for the message and produced deletes downstream the message key is removed from the filter table and the offset is committed.

Notable Variations —

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help ∨

Multiple Topic Filter

Filter steps can be configured with multiple types. Example Step Configuration

Each filtered type is used in downstream steps by configuring <code>output_type_prefix_type</code>. The above example has 3 output types, <code>filtered_phone</code>, <code>filtered_person</code>, <code>filtered_team</code>.

Lookup ⁽⁻⁾

Lookup steps (type: lookup) are used to store data from a single raw topic in the persistence store which can then be queried directly from the mapper. Lookup steps can be used on either raw types or filtered raw types. This is a replacement for a common pattern of using persist-only shared key steps and querying the data from the mapper.

Note Lookup steps cannot be referenced by other steps.

Changes to a record specified in the type of a lookup steps configuration will not trigger a remapping of the messages that previously used it. A full example of this initialization can be seen in the lookup-table example mapper.

Cache



Entity Builder About About Getting Started Cost Tour Architecture Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam ^

Getting Started

Deployments

Beam Operations

Help

For those utilizing the ebaas-mapper-common-golang framework, you can see example usage here. The lookup cache will keep records in a cache in the memory of the mapper. Canonical messages can then be enriched with this data. The cache can be configured using some of the parameters found here. Cache size configuration is based on the number of unique keys. If you're using a grouping key lookup, you should pay extra attention to this. The number of grouping keys will be limited by cache size, not the total number of rows.

Example Step Configuration

```
- step_name: state_lookup
                                       type: lookup
    lookup_config:
      type: state
      keys:
        stateCD
```

A full list of configuration options can be found in **Lookup Configuration**

How it works

- 1. Entity Joiner receives a message from a lookup type topic
- 2. Entity Joiner writes row to the Lookup table. This uses the persistence table, but with the table name lookup.
- 3. Offset for the message is committed and joiner moves on to the next message

Message deleted

- 1. Entity Joiner receives a message from a lookup type topic
- 2. Entity Joiner uses the key fields to delete the row from the Lookup table
- 3. Offset for the message is committed and joiner moves on to the next message

Transform Only -

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments \

Help ✓

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help

A transform only step (type: transform) offers a way to map a single raw record to a canonical message. By default, it is stateless, meaning no data is persisted during transformation. This step is handled entirely in the Entity Transformer and Mapper.

Note Because there is no state this step can only be used to map 1 raw record to a 1 canonical message.

Example Step Configuration

```
- step_name: member
    type: transform
    transform_config:
    type: raw_member
```

A full list of configuration options can be found in Transformation Configuration

How it works

- Entity Transformer receives a message from the raw topic
- 2. Sends message to the Mapper
- Mapper returns a canonical message and the Transformer produces to the output topic.

Pass Through ⁽⁻⁾

A Pass Through step (type: pass-through) allows you to use a code-based filter on a single topic, by providing a pass through container rather than a traditional mapper. It does not allow for the transformation of the messages themselves, it only allows for the filtering of which messages make it to the output topic.

Pass Through is generally transform only, and stateless. If a persist-only step, like lookup, is used then this is no longer the case, as the joiner and the persistence store will be utilized. It is also no longer stateless if minimize_tombstones is enabled.

One of the benefits of pass-through, as opposed to using transform, is that you do not need to maintain



About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations ~

Help

canonical schemas within the containers. The schema of the canonical messages will always be the same as those of the raw topics. This also means that if your raw topic schema changes slightly, your container filtering logic may continue to work, as long as the changes are not relevant to the workings of that logic.

Pass Through has its own contract with the transformer, and all Pass Through containers must implement this contract. Currently, EBaaS supports pass through contract containers in its golang mapper library. A full library-utilizing container example can be found here. Users can also implement the contract in the language of their choice.

Example Step Configuration

```
- step_name: member
    type: pass-through
    pass_through_config:
        type: raw_member
```

How it works

- 1. Entity Transformer receives a message from the raw topic
- 2. Sends message to the Pass Through container, where the message is filtered
- The container returns a status and the Transformer produces an unmodified message to the output topic based on this status.

Fan Out ⁽⁼⁾

The fan out step allows you to add fields from an array field in message value to the message key for joining purposes. Each element in the array will be sent as a message to downstream steps. An array with 5 elements will be sent to downstream steps 5 times each with a key containing the information pulled from that element.

Note A fan out step will only allow you to join by an array by pulling fields into the key. The original message value is sent to downstream steps unchanged.

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture

Topology Configuration

Topology Modification

Optimization Guide

Containers

Topology

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments \

Beam Operations

Help

Example Step Configuration

```
- step_name: fan_out_book
   type: fan-out
   fan_out_config:
     type: book
   keys:
     - bookID
   array_path: authors
   fan_out_keys:
     - path: authorID>int
     output_key: authorID
```

A full list of configuration options can be found in Fan Out Configuration

How it works

Create/Update

- 1. Entity Joiner receives a message from the child topic and builds.
- 2. Reads fan out table to read previous value keys for the message.
- 3. For each element in the array field to fan out by
 - i. A message is created with the fan out key fields added to the message key
 - ii. The modified message is sent to downstream steps to be processed
 - iii. If fan out keys were not in the fan out table previously, a new row is written to the fan out table
- 4. For each row in the fan out table that was not found in the array to fan out
 - i. A delete message is created with the fan out key fields added to the message key
 - ii. The delete message is sent to downstream steps to be processed
 - iii. The fan out table row is deleted

Delete

- 1. Entity Joiner receives a message from a topic
- 2. Reads fan out table to read previous value keys for the message.
- 3. For each row from the fan out table



- About
 About
 About
 Getting Started
 Cost
 Tour
 Architecture
 Topology
 - **Topology Configuration**
 - **Topology Modification**
 - **Optimization Guide**
 - Containers
 - Database
 - **Test Harness**
 - Template Mapper
 - Transformer
 - Deployments
 - Help
 - Terms of Service
- Beam
 - **Getting Started**
 - Deployments
 - Beam Operations
 - Help

- i. A delete message is created with the fan out key fields added to the message key
- ii. The delete message is sent to downstream steps to be processed
- iii. The fan out table row is deleted

Notable Variations

Fan Out by Row Type

Fan out steps can also be configured to fan out based on values of a non array field within a row type. A common use for this is fanning out a join only shared key step by values from a row type with multiple rows. To configure a case like this, array_path is not provided and row_type is required.

Example Step Configuration

```
- step_name: fan_out_book
   type: fan-out
   fan_out_config:
    type: book_join
   row_type: book
   keys:
        - bookID
   fan_out_keys:
        - path: authorID>int
        output_key: authorID
```

How it works

^

This configuration works identically to the fan out step with notable difference. While a fan out step on an array field does not modify the message value at all (the array remains in the value and only the key is modified), fanning out by a row type limits the rows sent to downstream steps for the row type to those associated to the fanned out key.

Internal Topics

Entity-Builder ensures that all messages for the same key are processed synchronously, in order, to prevent data issues due to race conditions. This is the reason for parent-child and join by value shared-key steps. To accomplish this, entity-builder uses several internal

Entity Builder About About Getting Started Cost Tour Architecture Topology **Topology Configuration**

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

topics in order to take advantage of the kafka partitions and the order guarantees they provide.

Join Queue (=)

The join queue is a topic used as the output topic for parent-child and join only shared-key steps.

The topic name is

ebaas_prefix | canonical_name | joinqueue optional encrypted suffix . ebaas_prefix is configured in topics.yaml canonical_name is configured as the name field in join.yaml. The topic will have a _encrypted suffix if beam encryption is enabled.

Ex.

EBAAS.hemi.attperf.order_payment_schedules.joi n-queue_encrypted

A single join queue is used for the entire entity-builder topology.

Message Format

Kev

field	description
stepName	The name of the step that produced the message. This is used to determine which step config to use to process the message.
parentKeyFields	This is an array of the field names/values from the parent_config of the step that produced it. For shared-key join by value this is populated with the shared keys.
childKeyFields	This is an array of the field names/values from the child_config of the step that

Entity Builder

About

About

Getting Started

Cost

Tour

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help ✓

field	description
	produced it. For shared-key join by value this is an empty array.

Example

```
{
   "stepName":"phone_w_person",
   "parentKeyFields":[{"name":"personID","val
   "childKeyFields":[{"name":"phoneID","value
}
```

Value ⁽⁼⁾

field	description
ChildType	The row_type for the child row in the step that produced the message.
ParentType	The row_type for the parent row in the step that produced the message.

Example

^

```
{
  "ChildType":"phone",
  "ParentType":"person"
}
```

Join message handling

Messages from the join queue are handled identically to a message that comes from a raw kafka topic, except that the "value" is read from the previous steps persistence table.

1. The stepName from the join queue message key as the previous step persistence table to read from.



- Entity Builder
 - About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help ∨

- 2. The parent key is constructed from the parentKeyFields
- 3. The previous parent row is read from the previous step persistence table where:
 - the partitioning_key is equal to the parent key
 - the row_type is equal to the ParentType from the message value
 - row_primary_keys is empty
- 4. The child key is constructed from the childKeyFields
- 5. The previous child row is read from the previous step persistence table where:
 - the partitioning_key is equal to the parent key
 - the row_type is equal to the ChildType from the message value
 - row_primary_keys is equal to the child key

The remaining joiner functionality is the same with the parent and child rows used as the "value" of the message.

Note Unlike a raw topic, the join queue message is never a tombstone. Instead, a message is treated as a delete if no child row is found in the previous step persistence table.

Publish Queue ⁽⁼⁾

The publish queue is a topic used by joiner to communicate to Entity Transformer that the data for shared-key has changed and needs to be remapped. Like the join queue, there is a single publish queue per topology.

The topic name is

ebaas_prefix . canonical_name .publish-queue optional encrypted suffix . ebaas_prefix is configured in topics.yaml . canonical_name is configured as the name field in join.yaml . The topic will have a _encrypted suffix if beam encryption is enabled.

About
About
About
Getting Started
Cost
Tour
Architecture

Topology Configuration

Topology Modification

Optimization Guide

Containers

Topology

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments ~

Beam Operations

Help

Ex.

EBAAS.hemi.attperf.order_payment_schedules.pub
lish-queue_encrypted

Message format —

Key

field	description
name	The shared-key step the message was produced by
key	The partitioning_key that needs to be remapped

Example

```
{
  "name":"team_canonical",
  "key":"2"
}
```

Value ⁽⁼⁾

field	description
data	An array of fields that make up the partitioning key. Used to construct the keyData map provided to the mapper

Example

```
{
  "data":[{"name":"teamID","value":"2","type
}
```

Low Priority Join Queue ⁽⁼⁾

The low priority join queue allows for prioritized joining. Message format is identical to the join queue.

The topic name is

ebaas_prefix . canonical_name .low-priority-join-

About
About
About
Getting Started
Cost
Tour
Architecture
Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

queue. ebaas_prefix is configured in topics.yaml. canonical_name is configured as the name field in join.yaml.

Ex.

```
EBAAS.hemi.attperf.order_payment_schedules.low -priority-join-queue
```

A single low priority join queue is used for the entire entity-builder topology.

Message handling —

Messages are read from low priority join topic after the configured time without a message from a higher priority topic (Consume Priority Configuration).

- 1. Message is read from the low priority join queue
- 2. (ebaas-priority) header is added to the message
- 3. Message is produced to the join queue.

Messages are redirected to the standard join queue to prevent race conditions that could arise with the same join key being processed from both the join queue and low priority join queue. The <code>ebaas-priority</code> header is used to ensure that messages from the low priority join queue result in messages produced back to the low priority join queue and the low priority publish queue by down stream steps.

Low Priority Publish Queue ⁽⁼⁾

The low priority publish queue allows for prioritized transformation. Message format is identical to the publish queue.

The topic name is

```
ebaas_prefix.canonical_name.low-priority-publish-queue.ebaas_prefix is configured in topics.yaml.canonical_name is configured as the name field in join.yaml.
```

Ex.

EBAAS.hemi.attperf.order_payment_schedules.low -priority-publish-queue



Cost

About
About
Getting Started

Tour
Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ✓

Terms of Service

Beam

Getting Started

Deployments \

Beam Operations

Help

A single low priority publish queue is used for the entire entity-builder topology.

Message handling —

Messages are read from low priority join topic after the configured time without a message from a higher priority topic (Consume Priority Configuration).

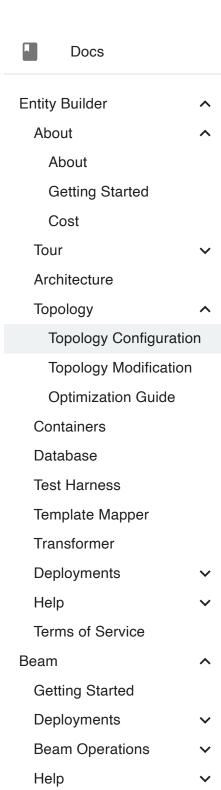
- 1. Message is read from the low priority publish queue
- 2. Message is produced to the publish queue.

Messages are redirected to the standard publish queue to prevent race conditions that could arise with the same publish key being transformed from both the publish queue and low priority publish queue.

Configuration Reference

General configs ⁽⁻⁾

Field	
name	The name of to construct of topic names
global_config	An optional f for the joiner
global_config.key_separator	You shouldn'
global_config.empty_key_substitute	You shouldn'
global_config.trim_front_whitespace	When set to whitespace for join. Applies key and mes
global_config.trim_back_whitespace	When set to whitespace find join. Applies key and mes



Field	
topology_steps	The list of joi
topology_steps.step_name	The name of the step. Res subsequents the type in parent_chi parent_chi or models.t
topology_steps.type	The type of t

Shared Key Configuration

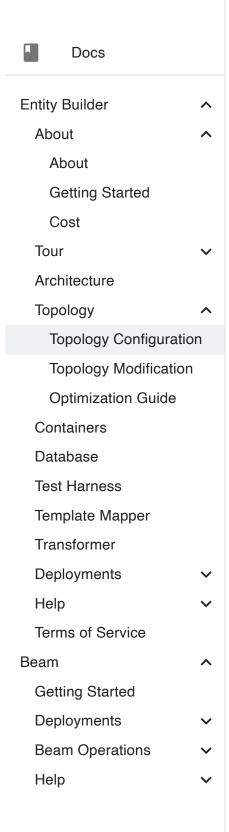
Field

topology_steps.minimize_tombstones

		Field
Docs		
Entity Builder	^	
About	^	
About		
Getting Started		
Cost		
Tour	~	
Architecture		
Topology	^	
Topology Configuration	n	
Topology Modification	ı	topology_steps.changed_field_tracking
Optimization Guide		
Containers		
Database		
Test Harness		
Template Mapper		
Transformer		
Deployments	~	
Help	~	
Terms of Service		
Beam	^	
Getting Started		
Deployments	~	topology_steps.changed_field_tracking_settings.arra
Beam Operations	~	
Help	~	
		topology_steps.changed_field_tracking_settings.arra

		Field
Docs		
Entity Builder	^	
About	^	
About		
Getting Started		
Cost		
Tour	~	topology_steps.changed_field_tracking_settings.arra
Architecture		
Topology	^	
Topology Configurat	ion	
Topology Modification	on	
Optimization Guide		
Containers		topology_steps.models
Database		
Test Harness		
Template Mapper		topology_steps.models.shared_keys
Transformer		
Deployments	~	
Help	~	topology_steps.models.topics
Terms of Service		
Beam	^	
Getting Started		
Deployments	~	
Beam Operations	~	topology_steps.models.topics.type
Help	~	
		topology_steps.models.topics.row_type_override

		Field
Docs		
Entity Builder	^	
About	^	
About		
Getting Started		topology_steps.models.topics.additional_keys
Cost		
Tour	~	
Architecture		
Topology	^	
Topology Configurati	on	
Topology Modificatio	n	topology_steps.models.topics.shared_key_overrides
Optimization Guide		
Containers		
Database		
Test Harness		topology_steps.models.topics.shared_key_overrides
Template Mapper		topology_ctopo.modolo.topico.cnarou_itoy_cvomdoc
Transformer		
Deployments	~	
Help	~	topology_steps.models.topics.shared_key_overrides
Terms of Service		
Beam	^	
Getting Started		
Deployments	~	
Beam Operations	~	topology_steps.models.topics.shared_key_overrides
Help	~	
		topology_steps.models.topics.shared_key_overrides



	Field
topology_steps.join_only	
topology_steps.persist_only	

Parent Child Configuration =

Field topology_steps.parent_child_configs topology_steps.parent_child_configs.buckets

Entity Builder About About **Getting Started** Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database Test Harness Template Mapper Transformer Deployments Help Terms of Service Beam **Getting Started** Deployments **Beam Operations** Help

Docs

topology_steps.parent_child_configs.child_config.type topology_steps.parent_child_configs.child_config.ro topology_steps.parent_child_configs.child_config.ke topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa		Fie	ld
topology_steps.parent_child_configs.child_config.ro topology_steps.parent_child_configs.child_config.ke topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config
topology_steps.parent_child_configs.child_config.ke topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.typ
topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.rov
topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.ke
topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.pa
topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.pa
topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.pa
topology_steps.parent_child_configs.child_config.pa	topology	_steps.parent_child_configs.c	child_config.pa
	topology	_steps.parent_child_configs.c	child_config.pa
topology_steps.parent_child_configs.child_config.pa	topology	steps.parent_child_configs.c	child_config.pa
	topology	_steps.parent_child_configs.c	child_config.pa

Docs **Entity Builder** About About **Getting Started** Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database Test Harness Template Mapper Transformer Deployments Help Terms of Service Beam **Getting Started** Deployments **Beam Operations** Help

Field
topology_steps.parent_child_configs.parent_config
topology_steps.parent_child_configs.parent_config.t
topology_steps.parent_child_configs.parent_config.re
topology_steps.parent_child_configs.parent_config.k
topology_steps.parent_child_configs.parent_config.a
topology_steps.parent_child_configs.persist_parent_

Rekey Configuration

Field
topology_steps.rekey_config
topology_steps.rekey_config.type
topology_steps.rekey_config.keys

		Field
Docs		
Entity Builder	^	
About	^	topology_steps.rekey_config.keys_from_value
About		
Getting Started		
Cost		
Tour	~	to a large state and the same first to the same state and the
Architecture		topology_steps.rekey_config.keys_from_value.path
Topology	^	
Topology Configurati	on	
Topology Modification		
Optimization Guide		
Containers		topology_steps.rekey_config.keys_from_value.outpu
Database		
Test Harness		
Template Mapper		
Transformer		
Deployments	~	
Help	~	topology_steps.rekey_config.keys_from_value.row_t
Terms of Service		
Beam	^	
Getting Started		
Deployments	~	
Beam Operations	~	topology_steps.rekey_config.keys_from_value.outpu
Help	~	

Filter Configuration

Field Docs topology_steps.filter_config **Entity Builder** topology_steps.filter_config.type About About **Getting Started** topology_steps.filter_config.types Cost Tour Architecture Topology topology_steps.filter_config.output_type_prefix **Topology Configuration Topology Modification** Optimization Guide Containers Database **Test Harness** Template Mapper topology_steps.filter_config.any_match Transformer **Deployments** Help Terms of Service Beam topology_steps.filter_config.any_match.conditions **Getting Started** Deployments **Beam Operations** topology_steps.filter_config.any_match.conditions.fie Help topology_steps.filter_config.any_match.conditions.frc

		Field
Docs		
Entity Builder	^	topology_steps.filter_config.any_match.conditions.wi
About	^	
About		
Getting Started		
Cost		
Tour	~	topology_steps.filter_config.any_match.conditions.bl
Architecture		
Topology	^	
Topology Configurati	on	
Topology Modificatio	n	
Optimization Guide		to a la sur atoma filtare a suffer a sur mantale a sur ditiona sul
Containers		topology_steps.filter_config.any_match.conditions.wl
Database		
Test Harness		
Template Mapper		
Transformer		
Deployments	~	
Help	~	topology_steps.filter_config.any_match.conditions.ch
Terms of Service		
Beam	^	
Getting Started		
Deployments	~	
Beam Operations	~	
Help	~	topology_steps.filter_config.any_match.conditions.ig
		topology_steps.filter_config.any_match.conditions.fie

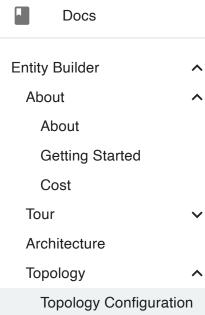
		Field
Docs		
Entity Builder	^	
About	^	topology_steps.filter_config.any_match.conditions.fie
About		
Getting Started		
Cost		
Tour	~	topology_steps.filter_config.any_match.conditions.fie
Architecture		
Topology	^	
Topology Configuration	on	
Topology Modification	า	
Optimization Guide		topology_steps.filter_config.any_match.conditions.nu
Containers		
Database		
Test Harness		
Template Mapper		
Transformer		topology_steps.filter_config.any_match.conditions.nu
Deployments	~	
Help	~	
Terms of Service		
Beam	^	
Getting Started		topology_steps.filter_config.any_match.conditions.nu
Deployments	~	
Beam Operations	~	
Help	~	
		topology_steps.filter_config.any_match.conditions.nu
		topology_steps.filter_config.any_match.conditions.nu

Docs **Entity Builder** About About **Getting Started** Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database **Test Harness** Template Mapper Transformer Deployments Help Terms of Service Beam **Getting Started** Deployments **Beam Operations** Help

	Field
topology_steps.filter_config.any_match.co	onditions.da
topology_steps.filter_config.any_match.co	onditions.da
topology_steps.filter_config.any_match.co	onditions.da

Lookup Configuration

Field	De
topology_steps.lookup_config	Hol cor for ste
topology_steps.lookup_config.type	The tha the use the cac
topology_steps.lookup_config.keys	A li fiel rec



Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

Field	De
	sav
	per
	sto
	que
	cac
	A li
	ren
	key
topology_steps.lookup_config.additional_keys	gro
	rec
	the
	loo

Transform Only Configuration ⁽⁼⁾

Field	Descriptio
topology_steps.minimize_tombstones	Allows the transformer to optimize tombstones by storing a additional row in the persistence store table that is checked after the
	mapping. If the last produce wa a delete, and the current mapper result is als a delete, transformer will not

Docs	
Entity Builder	^
About	^
About	
Getting Started	
Cost	
Tour	~
Architecture	
Topology	^
Topology Config	guration
Topology Modifi	cation
Optimization Gu	uide
Containers	
Database	
Test Harness	
Template Mapper	
Transformer	
Deployments	~
Help	~
Terms of Service	
Beam	^
Getting Started	
Deployments	~
Beam Operations	~

Help

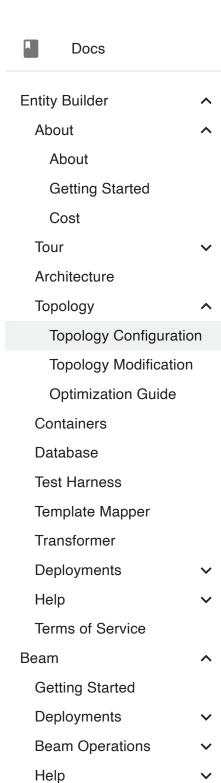
Field	Descriptio
	produce to the canonical topic. Usefu for topologies where records frequently cause canonical deletes. Makes transform stateful
topology_steps.transform_config	Holds configuratio for a transform step
topology_steps.transform_config.type	The type that maps to the topic to be transformed This type can be filtered

Pass Through Configuration

Field	Descri
topology_steps.minimize_tombstones	Allows transfor to optim tombsto by stori additior

Docs	
Entity Builder	^
About	^
About	
Getting Started	
Cost	
Tour	~
Architecture	
Topology	^
Topology Confi	guration
Topology Modif	cation
Optimization G	uide
Containers	
Database	
Test Harness	
Template Mapper	,
Transformer	
Deployments	~
Help	~
Terms of Service	
Beam	^
Getting Started	
Deployments	~
Beam Operations	· •
Help	~

Field	Descri
	row in t
	persiste
	store ta
	that is
	checke
	after the
	mappin
	the last
	produce
	a delete
	and the
	current
	mapper
	result is
	a delete
	transfor will not
	produce the
	canonic
	topic. U
	topolog
	where
	records
	frequen
	cause
	canonic
	deletes
	Makes
	through
	stateful
	Holds
topology_steps.pass_through_config	configu
. 3,2	for a pa
	through
topology_steps.pass_through_config.type	The typ
	that ma
	the topi

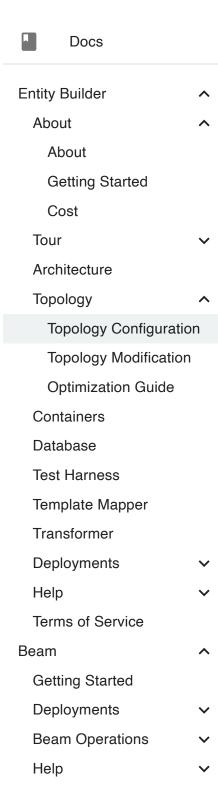


Field	Descri
	messaç
	are beir
	passed
	passed type ca filtered
	filtered

Fan Out Configuration ⁽⁼⁾

. un cut comiguration
Field
topology_steps.fan_out_config
topology_steps.fan_out_config.type
topology_steps.fan_out_config.keys
topology_steps.fan_out_config.array_path
topology_steps.fan_out_config.row_type

topology_steps.fan_out_config.fan_out_keys



Field
topology_steps.fan_out_config.fan_out_keys.path
topology_steps.fan_out_config.fan_out_keys.output_

Consume Priority environment variables

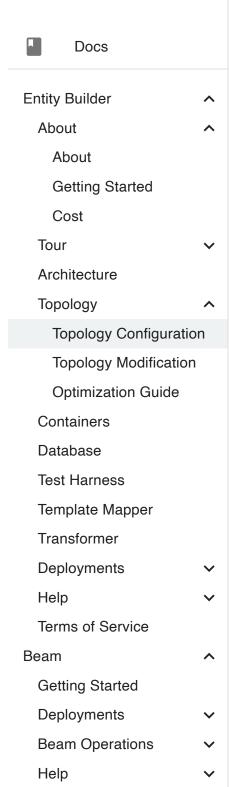
 Θ

There are several environment variables that can be used to tune consume priority

Env	De
CONSUME_PRIORITY_TIMEOUT_MS	The of mill tha pas with hig pric me bei low pric me will

Docs	
Entity Builder	^
About	^
About	
Getting Started	
Cost	
Tour	~
Architecture	
Topology	^
Topology Configuratio	n
Topology Modification	
Optimization Guide	
Containers	
Database	
Test Harness	
Template Mapper	
Transformer	
Deployments	~
Help	~
Terms of Service	
Beam	^
Getting Started	
Deployments	~
Beam Operations	~
Help	~

Env	De
CONSUME_PRIORITY_LOWEST_PRIORITY	Deflow cor price alw high price is to price pric
CONSUME_PRIORITY_MIN_WAIT_MS	The mir nur mill tha pas a m from par be App all excepting



Env	De
	pric
	(pri
	A v zer
	hav
	mir
	del
	The
	ma
	nur
	mil
	the
	der will
	witl
	rea
	me
	fror
CONSUME_PRIORITY_MAX_WAIT_MS	par
	Apı
	all
	exc
	hig pric
	(pri
	Αv
	zer
	hav
	ma
	del

Logical Decimal Join Formatting



WARNING

Because this changes the format of the join key, changing this setting requires a full reload.

By default, logical decimal avro fields are formatted in the rational format provided by the goavro library that

Docs

Entity Builder About About **Getting Started** Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database **Test Harness** Template Mapper Transformer **Deployments** Help Terms of Service Beam

Getting Started

Beam Operations

Deployments

Help

Entity Builder uses. For instance, the integer 3 is formatted to staring as 3/1 and the decimal 1.5 is formatted as 3/2. This does not work very well when joining to another field type. You can change this formatting to be a more standard numeric format by providing an env FORMAT_BYTES_DECIMAL_FOR_JOIN set to true. When enabled, the string format used to join will be based on the scale property of the bytes decimal field. The scale property determines the number of decimal places to include in the string.

Examples

Rational Format	scale	Joined format
3/1	0	3
3/1	1	3.0
3/1	2	3.00
3/1	3	3.000

	Descr
FORMAT_BYTES_DECIMAL_FOR_JOIN	Formate decimal into stri represe of the n value w joining

This feature can be turned on in the test harness by providing the --decimal-join=true flag when executing test.sh or by adding the above env to eb_env

Topics =

topics.yaml defines the relationship between types in join.yaml and Kafka topics. This allows for easier deployments to multiple data centers.

Docs

Entity Builder About About **Getting Started** Cost Tour Architecture Topology **Topology Configuration Topology Modification** Optimization Guide Containers Database **Test Harness** Template Mapper Transformer **Deployments** Help Terms of Service Beam ^ **Getting Started** Deployments **Beam Operations**

Help

join.yaml can remain the same and only the topics must change to match the environment. Users need to specify all raw input topics and all (non join_only) shared-key steps. The type must match what is in join.yaml exactly.

Example:

In this example, the topology consumes off of raw input topics

- DE_STREAMS.ctc.cdc.cdb.cnsm_cob_primacy.v
 2
- DE_STREAMS.ctc.cdc.cdb.cnsm_mdcr_primacy.
 v2
- DE_STREAMS.ctc.cdc.cdb.l_cnsm_srch.v2

And produces to

 cdb.prod-eligibilityctc.prod.canonical.eligibility.coordination-ofbenefits.v3

Topic Priority

Individual topics can be given a priority between 0 (highest priority) and CONSUME_PRIORITY_LOWEST_PRIORITY. A topic with a priority set to a value other than 0 will wait for a configured amount of time without processing messages with a higher priority before a message is read from it (Consume Priority configuration. Steps consuming messages from topics with lowest priority

Docs

Entity Builder

About

About

Getting Started

Cost

Tour ~

Architecture

Topology

Topology Configuration

Topology Modification

Optimization Guide

Containers

Database

Test Harness

Template Mapper

Transformer

Deployments

Help ∨

Terms of Service

Beam

Getting Started

Deployments

Beam Operations

Help

will produce to low priority join queue and low priority publish queue.

Example:

 $output_topic: cdb.prod-eligibility-ctc.p\red.$

ebaas_prefix: EBAAS.cdb

topic_configs:

- type: cnsm_cob_primacy

topic_name: DE_STREAMS.ctc.cdc.cdb.cnsm_

- type: cnsm_mdcr_primacy

topic_name: DE_STREAMS.ctc.cdc.cdb.cnsm_

priority: 1

- type: l_cnsm_srch

topic_name: DE_STREAMS.ctc.cdc.cdb.l_cns

priority: 2

- type: coordination_of_benefits

topic_name: cdb.prod-eligibility-ctc.pro

Topic configs ⁽⁻⁾

Field	Description	
output_topic	The output topic for this topology	required
ebaas_prefix	Prefix used for internal topics	required
topic_configs	A list of type to topic mappings	required, at I
topic_configs.type	The type for this topic. Used in join.yaml	required
topic_configs.topic	The topic that corresponds to this type	required



Entity Builder	^
About	^
About	
Getting Started	
Cost	
Tour	~
Architecture	
Topology	^
Topology Configuration	n
Topology Modification	
Optimization Guide	
Containers	
Database	
Test Harness	
Template Mapper	
Transformer	
Deployments	~
Help	~

Terms of Service

Getting Started

Beam Operations

Deployments

Beam

Help

Field	Description	
topic_configs.priority	The priority of this topic.	min=0, max=CONSI